

Agent Reach

Give your AI agent eyes on the entire internet. One CLI, zero API keys — it works with any agent that can run a shell command.



Contents

Why this exists	3
Install — one line	5
What it can reach	7
Playbooks — 4 ways to use it	10
Stay safe	15
Why star it	16

Why this exists

Ask most agents to check a tweet, pull a YouTube transcript, or read a Reddit thread and they stall — paid APIs, IP blocks, login walls, or a pile of raw HTML they can't parse.

The usual fix costs money or time: a scraping API with a monthly bill, or an afternoon lost wiring up yet another CLI for one platform. Agent Reach is the capability layer that already solved "which tool, which fallback, which auth" for every platform that matters. You don't pick a scraper — you ask in plain language, and your agent reaches.

Who this is for

Solo founders doing their own outreach, growth marketers researching competitors, anyone mining Reddit or reviews for product ideas — if you're running lean and a \$50/mo scraping API doesn't pencil out for a one-off pull, this is built for you.

The repo

github.com/Panniantong/agent-reach — free, open-source, MIT licensed, 80k+ stars, GitHub Trending #1.

Works with

Claude Code, OpenClaw, Cursor, Windsurf — any agent that can run a shell command. It installs itself; you never touch a scraper directly.

What you can do

- Read any public page or PDF
- Pull YouTube & Bilibili transcripts
- Search the whole web — free, semantic
- Scrape gated platforms (one-time login)
- Subscribe to any RSS feed

Everything runs locally in your terminal. No API keys to buy. The four playbooks starting on page 10 show exactly how.



Chapter 1

Install

One line. Paste to your agent.

You don't install it — your agent does. Paste this into Claude Code (or any agent). It reads the doc, installs the CLI, runs dependency checks, and drops a `SKILL.md` so it knows what to reach for. `PASTE INTO YOUR AGENT`

COMMAND • COPY ME

Install Agent Reach for me:

`https://raw.githubusercontent.com/Panniantong/agent-reach/main/docs/install.md`


Default install is **read-only-safe**: it checks your environment and reports what's missing without touching your system. Add `--system` only once you're fine letting it install dependencies and write config.

Update later with the same pattern — swap `install.md` for `update.md`.

If you'd rather run it yourself

Once the CLI is on your machine: `agent-reach install --env=auto` checks only · add `--system` to actually install · add `--dry-run` to preview with no changes made.





Chapter 2

What it reaches

Works with zero setup

Just ask in plain language. Your agent maps it to the right tool underneath — nothing to configure.

ASK YOUR AGENT

WHAT RUNS UNDERNEATH

“Summarize this link”

Jina Reader — reads any public webpage

“What's this YouTube video about”

yt-dlp — pulls the transcript

“Search Bilibili for X”

bili-cli — no login needed

“Search the whole web for X”

Exa semantic search (MCP) — free

“Subscribe me to this RSS feed”

feedparser

“What does this GitHub repo do”

the `gh` CLI


A one-time "help me configure X"


Twitter/X, Reddit, Facebook, Instagram, LinkedIn, XiaoHongShu (小红书), Bilibili subtitles, Boss直聘, and 小宇宙 podcast transcription all unlock after you tell your agent to configure them — it walks you through login / cookie setup step by step. No manual docs to hunt down.

Run `agent-reach doctor` any time to see exactly which backend each platform is using, and what — if anything — needs fixing.

How it stays working

Every platform has a **primary + fallback** chain of backends (e.g. Twitter: `twitter-cli` → `OpenCLI` → `bird`). When a platform changes its blocking rules or a tool goes unmaintained, the maintainer swaps the chain — you get the fix by re-running update, not by re-learning a new tool.





Chapter 3

Playbooks

Build a lead list from LinkedIn + X

Thirty to fifty real prospects who match your ICP, with the context to write a first line that isn't generic — not a purchased list.

Step 1 — connect the accounts (one-time)

Tell your agent "help me configure LinkedIn" and "help me configure Twitter." LinkedIn routes through mcp-server-linkedin (Jina Reader as fallback for public pages); Twitter through twitter-cli, with OpenCLI as backup. Use burner accounts — see Chapter 4.

Step 2 — ask in plain language

No commands to memorize. Describe the search and your agent maps it to the right backend underneath.

ASK YOUR AGENT

"Search LinkedIn for [job title] at [industry] companies in [city] — pull their profile links."

"Search Twitter for people who posted about [problem] in the last week."

What good output looks like

A table your agent builds from the reads — name, headline or bio, profile link, and the specific line that makes them a fit (what they posted, what their bio says). Not raw scraped text you still have to sort through.

Pro tip

LinkedIn's zero-config path (Jina Reader) already reads public profile and company pages without any login — run a quick pass there before you invest in full LinkedIn config for search.



Reverse-engineer a competitor's YouTube + Instagram

Map what's actually working for a competitor or a creator you admire — their top content and the pattern behind it — before you plan your next month.

Step 1 — YouTube needs nothing

Subtitle extraction and video search run zero-config through yt-dlp. Just ask.

Step 2 — Instagram needs a one-time login

Tell your agent "help me configure Instagram." It routes through OpenCLI, reusing your desktop Chrome's existing session — unlocks user search, profile, a user's recent posts, and Explore.

ASK YOUR AGENT

"What are [channel]'s 10 most-watched videos, and what do the transcripts have in common?"


"Pull [competitor]'s Instagram profile and their most recent posts."

What good output looks like

Not a list of videos — a pattern. Your agent should surface the repeated hook structure, topics, and posting cadence across the transcripts and captions, not just dump the raw text.

Pro tip

Ask for both platforms in the same read and compare cadence — YouTube long-form vs. Instagram short-form usually reveals which format they're actually betting on.



Mine Reddit for customer pain points

Real complaints in customers' own words that become content angles or a product roadmap — the actual language people use when they're frustrated, not a paraphrase.

Step 1 — configure once (Reddit has no free path)

Reddit's anonymous API is blocked platform-wide, so there's no zero-config option. Tell your agent "help me configure Reddit" — it sets up OpenCLI on desktop (reuses your browser's login) or rdt-cli with a cookie.

Step 2 — confirm it's healthy

Run `agent-reach doctor` once after configuring — it shows exactly which Reddit backend is active before you run a big pull.

ASK YOUR AGENT

```
"Search Reddit for people complaining about  
[your problem space], across r/[sub1] and  
r/[sub2]."
```

```
"Pull the top comments on this thread: [url]."
```

What good output looks like

Direct quotes with thread links attached, not summaries — the actual wording is what makes headlines and product copy land.

Pro tip

Ask for the pain point AND the workaround people are already using — the workaround is usually your feature roadmap.



Pull a creator's transcripts to study their hooks

You're not copying scripts — you're studying the first five seconds across ten-plus videos to find the hook formula a creator repeats.

Step 1 — YouTube, zero setup

yt-dlp pulls the full transcript for any public video. No login, no config, no waiting.

Step 2 — Bilibili, partial setup

Search and video details run zero-config via bili-cli. Subtitles need one extra step — tell your agent "help me configure Bilibili" to unlock OpenCLI for transcript pulls.

ASK YOUR AGENT

```
"Give me the full transcript of this video:  
[url]."
```

```
"Pull transcripts for [creator]'s last 10 videos  
and show me how each one opens."
```


What good output looks like

The opening two to three lines of each transcript, laid side by side — that's where the hook pattern shows up, not in the middle of the video.

Pro tip

Do this for your own last 10 posts too. Comparing your hooks against theirs, side by side, is the fastest way to see what you're leaving on the table.





Chapter 4

Stay safe

Read this before you connect an account

- Credentials and cookies live only in `~/.agent-reach/config.yaml` on your machine (file permissions 600) — never uploaded anywhere.
- **Use a throwaway account, not your main.** For any platform that needs login (Twitter, XiaoHongShu, Reddit, Facebook, Instagram), scripted access carries a real ban risk on the account it authenticates as.
- Full uninstall: `agent-reach uninstall` wipes `~/.agent-reach/`, the skill files, and MCP config. Add `--dry-run` to preview first, or `--keep-config` to drop only the skill files and keep your tokens for a reinstall.



Why star it

The maintainer runs this daily and actively tracks platform breakage — when a scraping path gets blocked, the fallback chain gets patched, and you pull the fix with the one-line update instead of debugging it yourself. A star keeps it on your radar and backs the person keeping it alive.

Get the repo + this guide

Follow **@kaiships.ai** and comment **SKILL** on the reel — the automation sends you the GitHub link and this PDF instantly.



kaiships.ai

FIELD GUIDE
AGENT REACH